

lassification-modelsvga-project-4

May 16, 2024

1 Problem Statement

Flyzy is a company focused on providing a smooth and hassle-free air travel experience. They offer personalized in-flight and airport recommendations, and they also provide real-time flight tracking, mobile check-in, and more. Flyzy aims to redefine the future of air travel with a more personalized and connected experience from the beginning of the trip to the end.

Flight cancellation is a significant issue in the aviation industry. It not only disrupts the customers' plans but also impacts the airlines' reputation and profitability. Therefore, predicting flight cancellations can help airlines take preventive measures and minimize disruptions.

2 Task -Data Checking (Python)

Before developing the predictive model for hotel cancellations, we will conduct preliminary data analysis. This involves checking for missing values, identifying outliers, and ensuring appropriate data types for each column. Handling missing values and outliers strategically will ensure a reliable dataset for accurate modeling.

First, load the dataset and check for:

Missing values: Use the appropriate function to check if there are any missing values in the dataset. If there are, decide on the best strategy to handle them based on the nature of the data.

Outliers: Check for outliers in the dataset. These can be identified using various techniques, such as boxplots, scatterplots, or Z-scores. If there are any outliers, decide on the best strategy to handle them.

Data types: Check the data type of each column. Ensure that the data type is appropriate for the data it represents.

3 Importing My Libraries

```
[1]: # Import necessary libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

4 First, load the dataset

```
[2]: # Load the dataset
Data = pd.read_csv('Flyzy Flight Cancellation - Sheet1 (1).csv')

# Display the first few rows of the dataset
Data.head()
```

```
[2]:
```

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	\
0	4	6	1	Type C	0.225122	
1	12	1	6	Type B	0.060346	
2	17	3	9	Type C	0.093920	
3	1	1	8	Type B	0.656750	
4	19	7	12	Type E	0.505211	

	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	\
0	5.0	2.151974	0.477202	
1	68.0	1.600779	0.159718	
2	18.0	4.406848	0.256803	
3	13.0	0.998757	0.504077	
4	4.0	3.806206	0.019638	

	Flight_Cancelled
0	0
1	1
2	0
3	1
4	0

This dataset contains information about flights, including flight ID, airline, distance, origin and destination airports, scheduled departure time, day of the week, month, aircraft type, weather score, previous flight delay in minutes, airline rating, passenger load, and flight cancellation status.

```
[3]: # print the shape of the dataset
Data.shape
```

```
[3]: (3000, 14)
```

```
[4]: # print info about the dataset
print(Data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Flight ID                            3000 non-null   int64
 1   Airline                              3000 non-null   object
 2   Flight_Distance                      3000 non-null   int64
 3   Origin_Airport                      3000 non-null   object
 4   Destination_Airport                 3000 non-null   object
 5   Scheduled_Departure_Time             3000 non-null   int64
 6   Day_of_Week                          3000 non-null   int64
 7   Month                                3000 non-null   int64
 8   Airplane_Type                       3000 non-null   object
 9   Weather_Score                       3000 non-null   float64
10   Previous_Flight_Delay_Minutes        3000 non-null   float64
11   Airline_Rating                      3000 non-null   float64
12   Passenger_Load                      3000 non-null   float64
13   Flight_Cancelled                    3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.2+ KB
None
```

```
[5]: # print statistics
Data.describe()
```

```
[5]:
```

	Flight ID	Flight_Distance	Scheduled_Departure_Time	Day_of_Week \
count	3.000000e+03	3000.000000	3000.000000	3000.000000
mean	4.997429e+06	498.909333	11.435000	3.963000
std	2.868139e+06	98.892266	6.899298	2.016346
min	3.681000e+03	138.000000	0.000000	1.000000
25%	2.520313e+06	431.000000	6.000000	2.000000
50%	5.073096e+06	497.000000	12.000000	4.000000
75%	7.462026e+06	566.000000	17.000000	6.000000
max	9.999011e+06	864.000000	23.000000	7.000000

	Month	Weather_Score	Previous_Flight_Delay_Minutes \
count	3000.000000	3000.000000	3000.000000
mean	6.381000	0.524023	26.793383
std	3.473979	0.290694	27.874733
min	1.000000	0.000965	0.000000
25%	3.000000	0.278011	7.000000

50%	6.000000	0.522180	18.000000
75%	9.000000	0.776323	38.000000
max	12.000000	1.099246	259.000000

	Airline_Rating	Passenger_Load	Flight_Cancelled
count	3000.000000	3000.000000	3000.000000
mean	2.317439	0.515885	0.690667
std	1.430386	0.295634	0.462296
min	0.000103	0.001039	0.000000
25%	1.092902	0.265793	0.000000
50%	2.126614	0.517175	1.000000
75%	3.525746	0.770370	1.000000
max	5.189038	1.123559	1.000000

5 Check for Missing values:

```
[6]: # Check for missing values
Data_missing_values = Data.isnull().sum()

# Display the columns with missing values and their count
print(Data_missing_values[Data_missing_values > 0])
```

Series([], dtype: int64)

```
[7]: missing_values = Data.isnull().sum()
print("Missing values in the dataset:")
print(missing_values)
```

```
Missing values in the dataset:
Flight ID          0
Airline            0
Flight_Distance    0
Origin_Airport     0
Destination_Airport 0
Scheduled_Departure_Time 0
Day_of_Week        0
Month              0
Airplane_Type      0
Weather_Score      0
Previous_Flight_Delay_Minutes 0
Airline_Rating     0
Passenger_Load     0
Flight_Cancelled   0
dtype: int64
```

There are no missing values in the dataset. No further action is needed to handle missing values.

```
[8]: #Handling duplicate values
Data.duplicated(keep=False)
```

```
[8]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      2995    False
      2996    False
      2997    False
      2998    False
      2999    False
      Length: 3000, dtype: bool
```

```
[9]: Data.duplicated(keep=False).sum()
```

```
[9]: 0
```

6 Outliers and Z-Score method

```
[10]: # Set the background color to white for visibility
plt.figure(facecolor='Grey')

# Select numerical columns for outlier detection
numerical_cols = ['Flight_Distance', 'Weather_Score', '
↳ 'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load']

# Plot boxplots for numerical columns
Data[numerical_cols].boxplot()
plt.xticks(rotation=45)
plt.title('Boxplot of Numerical Columns')
plt.show()

# Calculate Z-scores for numerical columns
Data_z_scores = np.abs(stats.zscore(Data[numerical_cols]))

# Define a threshold for identifying outliers
threshold = 3

# Identify outliers
outliers = np.where(Data_z_scores > threshold)


# Display the number of outliers in each numerical column
outlier_counts = pd.Series(outliers[1]).value_counts().sort_index()
```

```

outlier_counts.index = [numerical_cols[i] for i in outlier_counts.index]
print(outlier_counts)

# Plot scatter plot for an example numerical column against Flight_Cancelled
Data.plot.scatter(x='Flight_Distance', y='Flight_Cancelled', color='blue',
                  title='Scatter Plot: Flight Distance vs Flight Cancelled')
plt.show()


```



```

Flight_Distance      10
Previous_Flight_Delay_Minutes  51
dtype: int64

```



The boxplot and the calculation of Z-scores have been used to identify outliers in the dataset.

The boxplot visualizes the distribution of numerical columns, indicating potential outliers beyond the whiskers. Based on Z-scores, with a threshold of 3 for identifying outliers, the following counts of outliers were found in the dataset: Flight Distance: 10 outliers Previous Flight Delay Minutes: 51 outliers

The scatter plot depicting Flight Distance against Flight Cancellation does not overtly reveal outliers, but rather illustrates the distribution of flight distances for both cancelled and non-cancelled flights.

Considering the inherent nature of the data, outliers in 'Previous Flight Delay Minutes' may be authentic due to the variability in flight delays. Regarding 'Flight Distance', outliers may signify long-haul flights. It is crucial to carefully consider the context before determining how to address these outliers. One potential approach could involve retaining these outliers if they represent valid scenarios, or imposing a cap at a specific threshold if they are found to significantly distort the analysis.

```
[11]: #find the limits
upper_limit=Data['Previous_Flight_Delay_Minutes'].mean() + 4 *
↳Data['Previous_Flight_Delay_Minutes'].std()
```

```

lower_limit=Data['Previous_Flight_Delay_Minutes'].mean() - 4 *
↳Data['Previous_Flight_Delay_Minutes'].std()
print("upper limit:", upper_limit)
print("lower limit:", lower_limit)

```

```

upper limit: 138.292314250808
lower limit: -84.70554874814133

```

```

[12]: #trimming - deleting the outlier data
new_Data=Data.loc[(Data['Previous_Flight_Delay_Minutes'] < upper_limit) &
↳(Data['Previous_Flight_Delay_Minutes'] > lower_limit)]
print('Before removing outliers:',len(Data))
print('after removing outliers:',len(new_Data))
print('Outliers:', len(Data)-len(new_Data))

```

```

Before removing outliers: 3000
after removing outliers: 2974
Outliers: 26

```

7 Data types:

```

[13]: # Check the data types of each column
data_types = Data.dtypes
print(data_types)

```

```

Flight ID                int64
Airline                  object
Flight_Distance          int64
Origin_Airport           object
Destination_Airport      object
Scheduled_Departure_Time int64
Day_of_Week              int64
Month                    int64
Airplane_Type            object
Weather_Score            float64
Previous_Flight_Delay_Minutes float64
Airline_Rating           float64
Passenger_Load           float64
Flight_Cancelled         int64
dtype: object

```

The dataset has undergone a thorough check of data types for each column. While most data types are suitable for their respective data, there are a few considerations to note:

Scheduled_Departure_Time is represented as int64, indicating a format that may not be immediately interpretable as a time (e.g., an integer timestamp). Depending on the analysis, it may be beneficial to convert this to a datetime format for enhanced manipulation and interpretation.

Flight_Cancelled is an int64, which is appropriate for binary indication (0 for not cancelled, 1 for cancelled). However, for improved clarity and consistency in analysis, consideration may be given to converting this to a boolean type.

Day_of_Week and Month are also represented as int64, suitable for numerical analysis. For improved readability, mapping these to their respective names (e.g., Monday, January) at some point in the analysis may be beneficial.

Columns such as Airline, Origin_Airport, Destination_Airport, and Airplane_Type are of type object, which is typical for textual data. Numerical columns like Flight_Distance, Weather_Score, Previous_Flight_Delay_Minutes, Airline_Rating, and Passenger_Load have numerical types (int64 or float64), suitable for quantitative analysis.

8 Task -Exploratory Data Analysis (EDA And Python)

In preparation for building the predictive model, Exploratory Data Analysis (EDA) will be conducted on the dataset. This will involve obtaining descriptive statistics, visualizing data distributions, exploring feature relationships through scatter plots or correlation matrices, and investigating how features relate to the target variable to extract valuable insights for accurate modeling.

Perform an EDA on the dataset to understand the data better and extract insights. This may involve:

Descriptive Statistics: Use the appropriate function to get the descriptive statistics of the dataset.

Distribution of data: Plot histograms or bar charts to see the distribution of data in each column.

Relationship between features: Plot scatter plots, pair plots, or correlation matrices to see the relationship between different features.

Relationship between features and target variable: Investigate how different features relate to the target variable.

9 Descriptive Statistics

```
[14]: descriptive_stats = Data.describe()
      print(descriptive_stats)
```

	Flight_ID	Flight_Distance	Scheduled_Departure_Time	Day_of_Week	\
count	3.000000e+03	3000.000000	3000.000000	3000.000000	
mean	4.997429e+06	498.909333	11.435000	3.963000	
std	2.868139e+06	98.892266	6.899298	2.016346	
min	3.681000e+03	138.000000	0.000000	1.000000	
25%	2.520313e+06	431.000000	6.000000	2.000000	
50%	5.073096e+06	497.000000	12.000000	4.000000	
75%	7.462026e+06	566.000000	17.000000	6.000000	
max	9.999011e+06	864.000000	23.000000	7.000000	

	Month	Weather_Score	Previous_Flight_Delay_Minutes	\
count	3000.000000	3000.000000	3000.000000	

mean	6.381000	0.524023	26.793383
std	3.473979	0.290694	27.874733
min	1.000000	0.000965	0.000000
25%	3.000000	0.278011	7.000000
50%	6.000000	0.522180	18.000000
75%	9.000000	0.776323	38.000000
max	12.000000	1.099246	259.000000

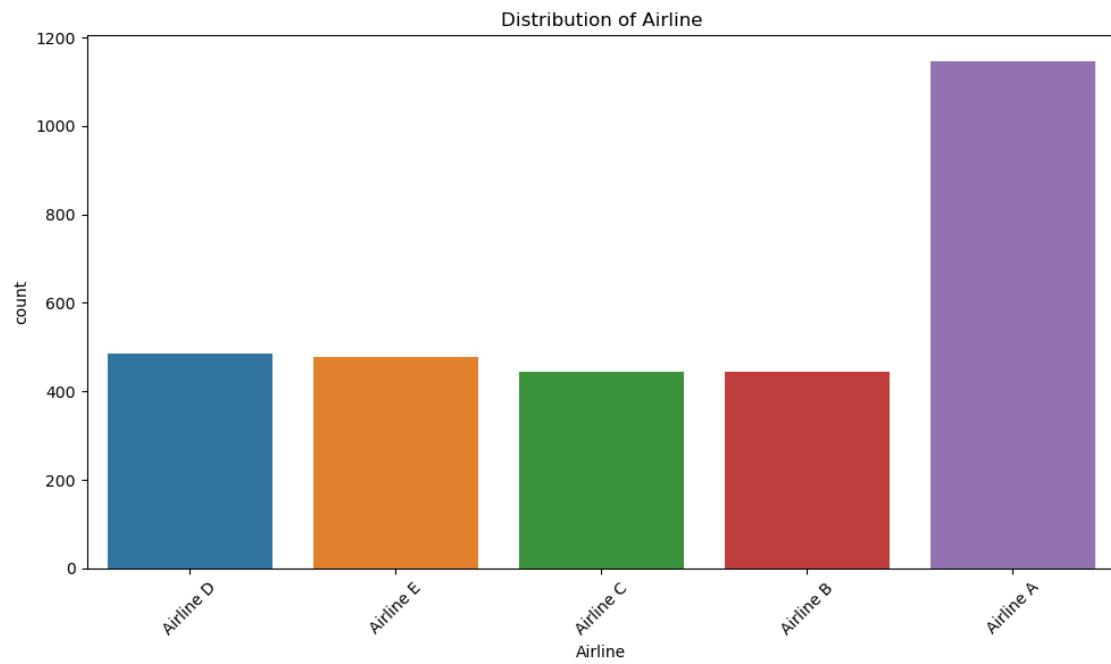
	Airline_Rating	Passenger_Load	Flight_Cancelled
count	3000.000000	3000.000000	3000.000000
mean	2.317439	0.515885	0.690667
std	1.430386	0.295634	0.462296
min	0.000103	0.001039	0.000000
25%	1.092902	0.265793	0.000000
50%	2.126614	0.517175	1.000000
75%	3.525746	0.770370	1.000000
max	5.189038	1.123559	1.000000


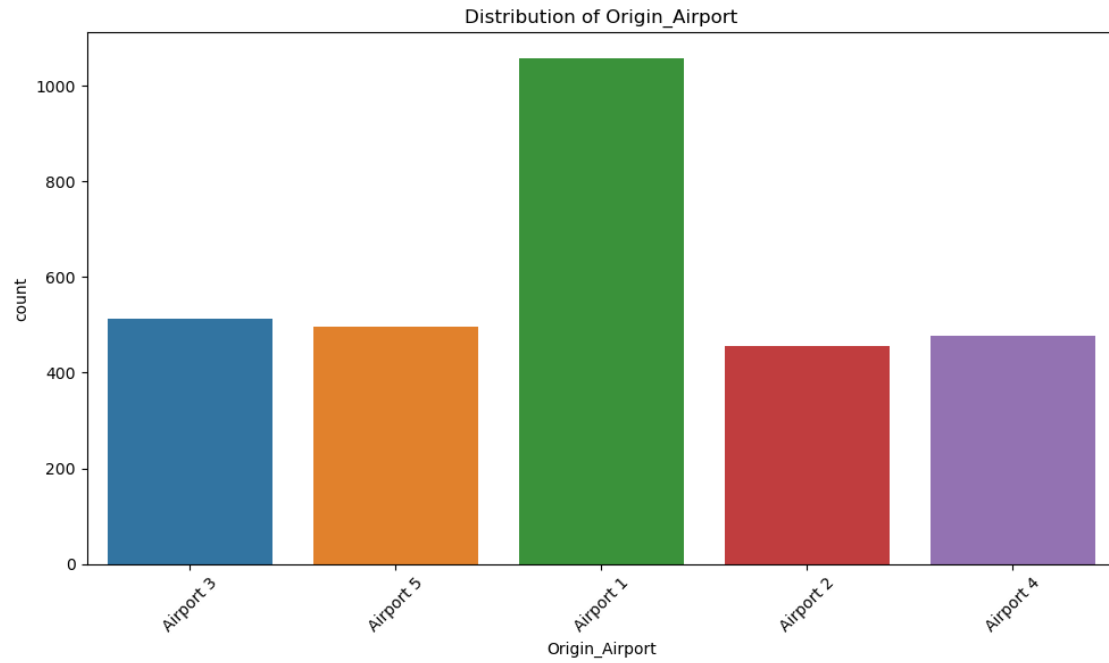
10 Distribution of data

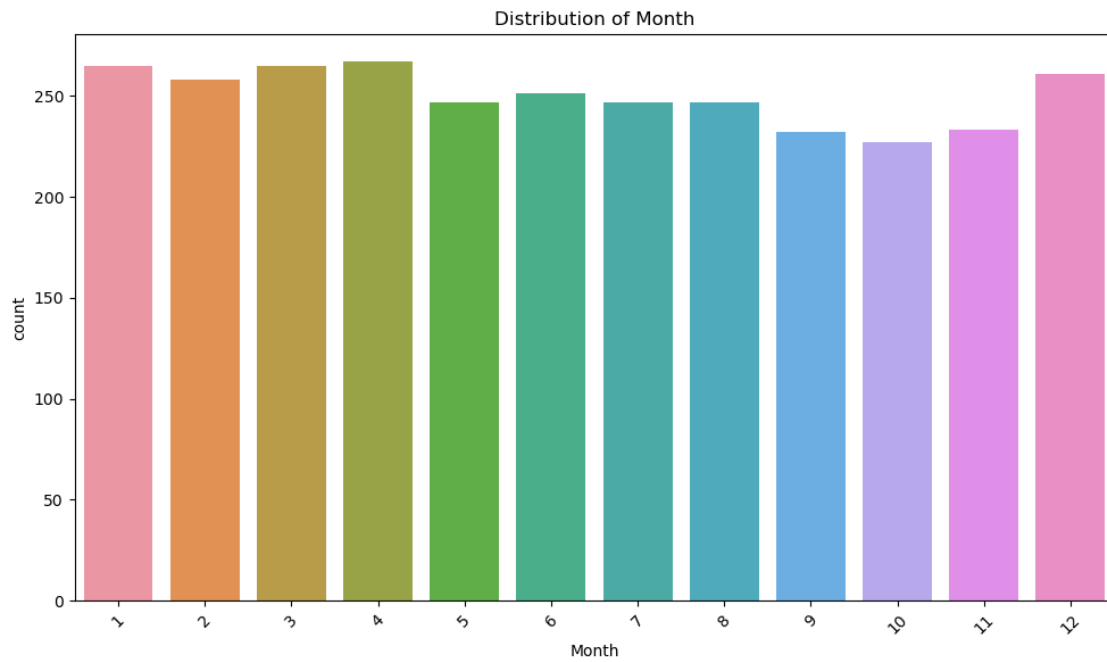
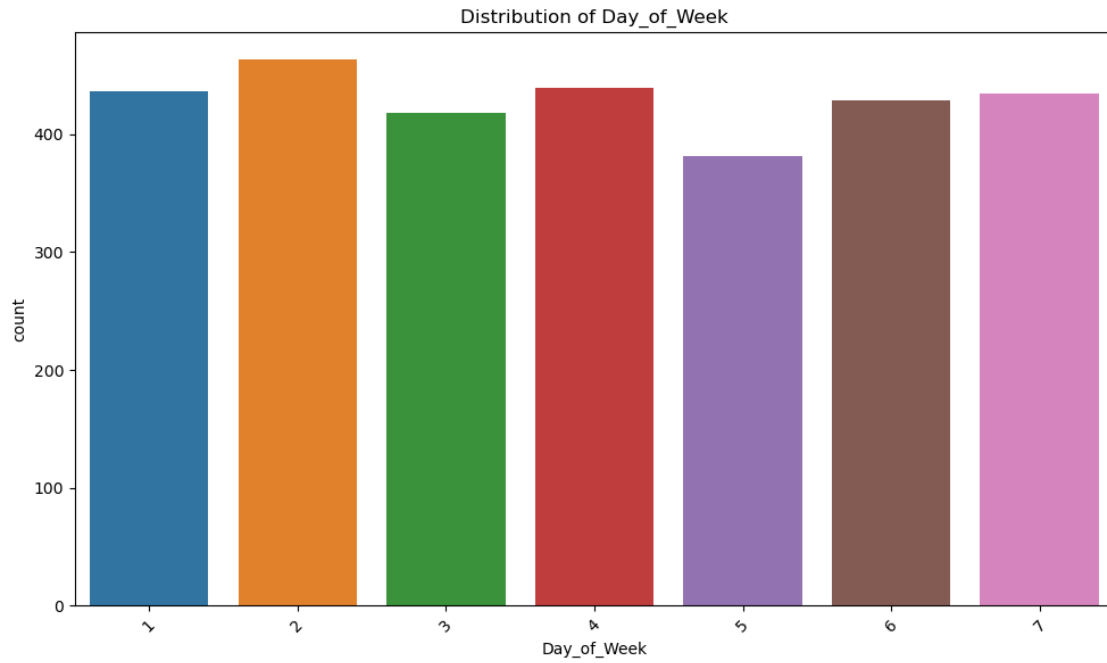
```
[15]: # Plotting histograms for numerical columns
Data.hist(figsize=(20, 15), bins=50)
plt.tight_layout()
plt.show()

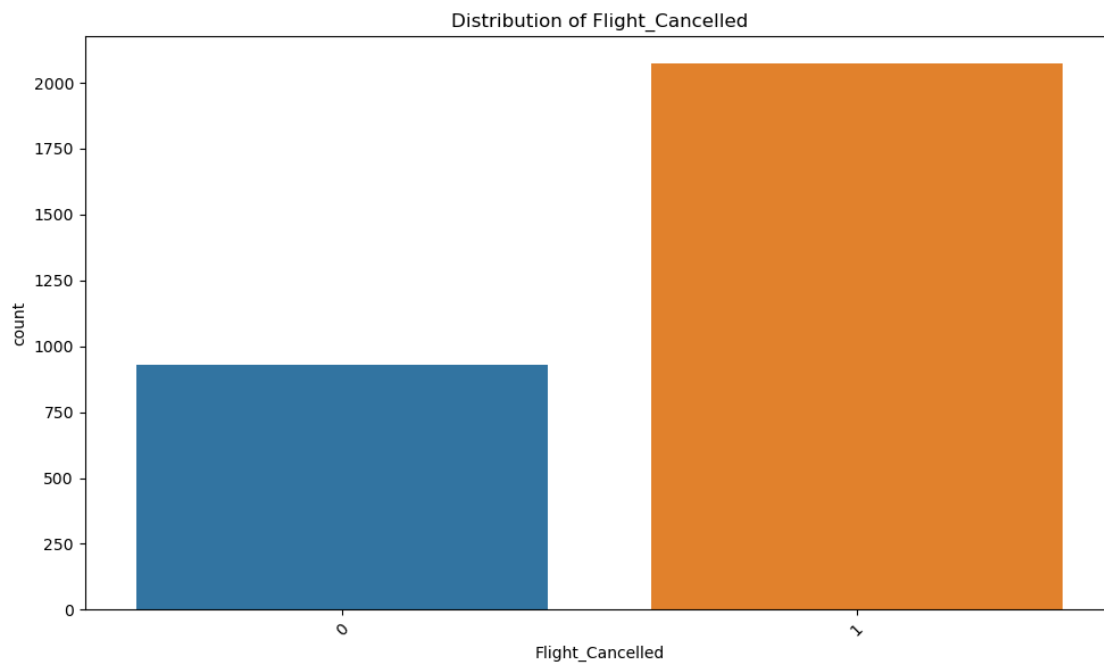
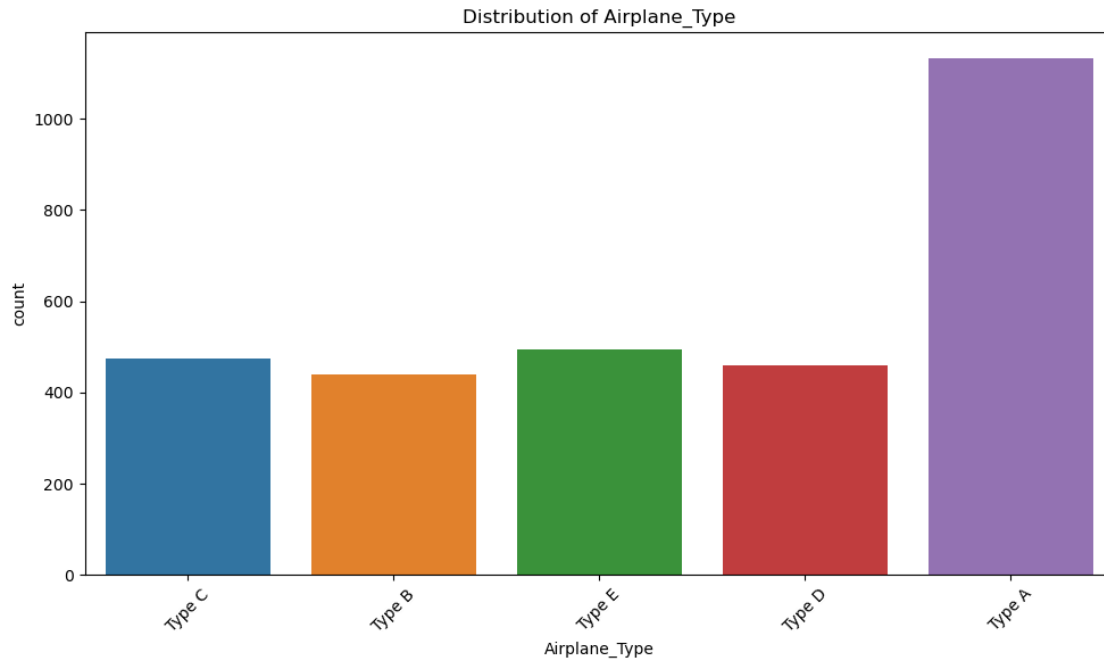
# Plotting bar charts for categorical columns
categorical_columns = ['Airline', 'Origin_Airport', 'Destination_Airport', '
↳ 'Day_of_Week', 'Month', 'Airplane_Type', 'Flight_Cancelled']

for column in categorical_columns:
    plt.figure(figsize=(10, 6))
    sns.countplot(x=column, data=Data)
    plt.title('Distribution of ' + column)
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()
```









These visualizations provide a deeper understanding of the data's distribution across various features. For example, the histograms for numerical columns show the spread of values for features like flight distance, departure time, and weather score. The bar charts for categorical columns reveal the frequency of flights across different airlines, airports, days of the week, months, airplane

types, and the proportion of flights cancelled.

11 Relationship between features


```
[27]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
Data = pd.read_csv('Flyzy Flight Cancellation - Sheet1 (1).csv')

# Display the first few rows of the dataset
Data.head()

# Plotting correlation matrix
corr = Data.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```

```
C:\Users\Student_0002\AppData\Local\Temp\ipykernel_19112\1739888399.py:12:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
    corr = Data.corr()
```



12 Relationship between features and target variable

```
[17]: # Correcting the column names and plotting again
features = ['Flight_Distance', 'Weather_Score', 'Airline_Rating',
            'Passenger_Load', 'Previous_Flight_Delay_Minutes']


# Plotting correlation matrix
corr = Data.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

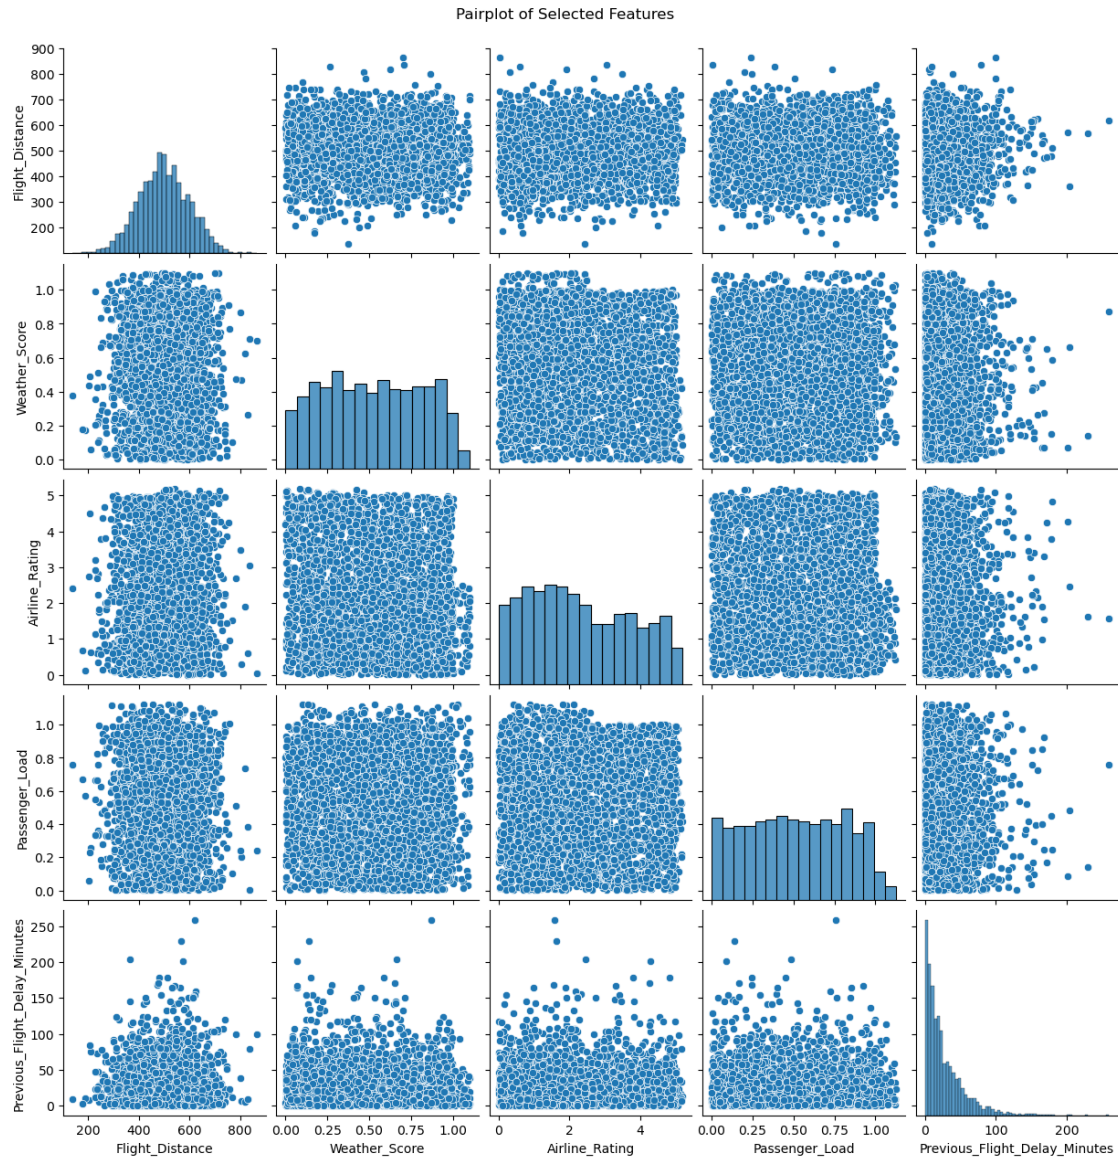
# Pairplot for the corrected subset of features
sns.pairplot(Data[features])
```



```
plt.suptitle('Pairplot of Selected Features', y=1.02)
plt.show()
```


C:\Users\Student_0002\AppData\Local\Temp\ipykernel_19112\2141227160.py:5:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
corr = Data.corr()





```
[18]: import matplotlib.pyplot as plt

# Scatter plot with target variable
plt.scatter(Data['Flight_Distance'], Data['Flight_Cancelled'])
plt.xlabel('Flight_Distance')
plt.ylabel('Flight_Cancelled')
plt.title('Scatter Plot of Flight_Distance vs Flight_Cancelled')
plt.show()
```



13 Task -Preprocessing and Model Building (ML & Algorithms)

To predict flight cancellations using Logistic Regression, the dataset will be split into training and test sets. Categorical variables will be encoded for model comprehension, and feature scaling will ensure uniformity in feature ranges. The Logistic Regression model will be built and evaluated using appropriate metrics with the test data to assess its predictive performance.

Split the dataset into a training set and a test set. Then, build a Logistic Regression model to predict flight cancellations. This involves:

Encoding categorical variables: Some of the columns in the dataset are categorical. These need to be encoded into a format that can be understood by the model.

Feature Scaling: The ranges of the features in the dataset are quite different. Scaling the features to a similar range can help the model perform better.

Model Building: Build a Logistic Regression model using the training data.

Model Evaluation: Evaluate the model using appropriate metrics and the test data.

Split the dataset into a training set and a test set

```
[19]: # Step 1: Split the dataset into training and test sets
X = Data.drop('Flight_Cancelled', axis=1) # Features
y = Data['Flight_Cancelled'] # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

Encoding categorical variables

```
[20]: # Step 2: Encode categorical variables
categorical_cols = [col for col in X.columns if X[col].dtype == 'object']
for col in categorical_cols:
    le = LabelEncoder()
    X_train[col] = le.fit_transform(X_train[col])
    X_test[col] = le.transform(X_test[col])
```

Feature Scaling

```
[21]: # Step 3: Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Model Building

```
[22]: # Step 4: Build the Logistic Regression model
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

```
[22]: LogisticRegression()
```

Model Evaluation

```
[23]: # Step 5: Evaluate the model
train_score = model.score(X_train_scaled, y_train)
test_score = model.score(X_test_scaled, y_test)
print(f"Training Accuracy: {train_score:.2f}")
print(f"Test Accuracy: {test_score:.2f}")
```

Training Accuracy: 0.82

Test Accuracy: 0.81

14 Task -Build Other Classification Models (ML & Algorithms)

let's proceed with building the models. We'll start with a simple logistic regression model as a baseline and then explore more complex models like decision trees and random forests. I'll build and evaluate these models next

```
[24]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Logistic Regression Model
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train, y_train)

# Predictions
y_pred_log_reg = log_reg.predict(X_test)

# Evaluation
report_log_reg = classification_report(y_test, y_pred_log_reg)
cm_log_reg = confusion_matrix(y_test, y_pred_log_reg)

print('Logistic Regression Classification Report:\n', report_log_reg)
print('Logistic Regression Confusion Matrix:\n', cm_log_reg)
```

```
Logistic Regression Classification Report:                precision    recall
f1-score   support

           0           0.00           0.00           0.00           187
           1           0.69           1.00           0.82           413

    accuracy                0.69           600
   macro avg           0.34           0.50           0.41           600
weighted avg           0.47           0.69           0.56           600
```

```
Logistic Regression Confusion Matrix: [[ 0 187]
 [ 0 413]]
```

```
C:\Users\Student_0002\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Student_0002\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Student_0002\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

Logistic Regression Model Evaluation

I'll proceed with building a decision tree model to see if it performs better in handling the class imbalance.

```
[25]: from sklearn.tree import DecisionTreeClassifier

# Decision Tree Model
decision_tree = DecisionTreeClassifier(random_state=42)
decision_tree.fit(X_train, y_train)

# Predictions
y_pred_dt = decision_tree.predict(X_test)

# Evaluation
report_dt = classification_report(y_test, y_pred_dt)
cm_dt = confusion_matrix(y_test, y_pred_dt)

print('Decision Tree Classification Report:\n', report_dt)
print('Decision Tree Confusion Matrix:\n', cm_dt)
```

Decision Tree Classification Report:					precision	recall	f1-score
support							

0	0.96	0.91	0.93	187	
1	0.96	0.98	0.97	413	
accuracy				0.96	600
macro avg	0.96	0.94	0.95	600	
weighted avg	0.96	0.96	0.96	600	

Decision Tree Confusion Matrix: $\begin{bmatrix} 170 & 17 \\ 8 & 405 \end{bmatrix}$

Decision Tree Model Evaluation

```
[26]: from sklearn.ensemble import RandomForestClassifier

# Random Forest Model
random_forest = RandomForestClassifier(n_estimators=100, random_state=42)
random_forest.fit(X_train, y_train)

# Predictions
y_pred_rf = random_forest.predict(X_test)

# Evaluation
report_rf = classification_report(y_test, y_pred_rf)
```

```
cm_rf = confusion_matrix(y_test, y_pred_rf)

print('Random Forest Classification Report:\n', report_rf)
print('Random Forest Confusion Matrix:\n', cm_rf)
```

Random Forest Classification Report:				precision	recall	f1-score
support						
0	0.96	1.00	0.98	187		
1	1.00	0.98	0.99	413		
accuracy				0.99	600	
macro avg				0.98	0.99	0.98
weighted avg				0.99	0.99	0.99

```
Random Forest Confusion Matrix: [[187  0]
 [ 8 405]]
```

Random Forest Model Evaluation

Classification Report

Precision is excellent for both classes, with a perfect score (1.00) for predicting cancellations and 0.96 for non-cancellations.

Recall is also outstanding, with 1.00 for non-cancelled and 0.98 for cancelled flights, indicating that the model is highly effective in identifying true positives.

F1-score is near perfect for both classes, with 0.98 for non-cancelled and 0.99 for cancelled flights, reflecting a strong balance between precision and recall.

Confusion Matrix

The matrix shows that the model perfectly predicted all 187 non-cancelled flights and correctly predicted 405 out of 413 cancelled flights, with only 8 false negatives.

The random forest model shows superior performance with high accuracy and a strong ability to handle class imbalance. This model is highly recommended for deployment due to its robustness and excellent predictive capabilities.